

Problem-1 (Assignment-1): Cache Hierarchy Optimization using gem5

Overview

This assignment requires you to investigate how different cache configurations affect processor performance using gem5 simulation. You will run a memory-intensive benchmark (matrix multiplication) with various cache parameters and analyze the results.

Learning Objectives

By completing this assignment, you will:

1. Understand cache hierarchy design and its impact on performance
2. Learn to use gem5 simulator for architecture exploration
3. Analyze performance metrics (execution time, hit rates, misses)
4. Make data-driven design decisions based on simulation results
5. Understand trade-offs between cache size, associativity, and cost

Assignment Tasks

Part 1: Environment Setup (10 minutes)

****Objective:**** Get familiar with the assignment tools

****Tasks:****

1. Verify gem5 RISCV build: `build/RISCV/gem5.opt --version`
2. Write the gem5 configuration script: `configs/cache_config.py`
3. Test single run with default cache configuration:

```
``` ./build/RISCV/gem5.opt configs/cache_config.py \
--l1i_size=16kB --l1d_size=16kB --l2_size=256kB \
--l1_assoc=4 --l2_assoc=8 \
--binary=<path_to_compiled_binary>
```

**\*\*Deliverables:\*\***

- Cache configuration script
- Screenshot of successful test run
- Output log showing cache statistics

#### ### Part 2: Single Parameter Sweep (30 minutes)

**\*\*Objective:\*\*** Understand impact of one cache parameter

**\*\*Choose ONE parameter to investigate:\*\***

- L1D cache size (vary: 16kB, 32kB, 64kB)
- L2 cache size (vary: 128kB, 256kB, 512kB, 1MB)
- L1 associativity (vary: 2, 4, 8)
- L2 associativity (vary: 4, 8, 16)

**\*\*Tasks:\*\***

1. Create a custom sweep script (`custom\_sweep.py`) for your chosen parameter
2. Run simulations with different values while keeping other parameters constant
3. Collect execution time (ticks) and cache hit rates
4. Create a table and plot showing the relationship

**\*\*Sample Default Configuration:\*\***

- L1I: 16kB, assoc=4
- L1D: 16kB, assoc=4
- L2: 256kB, assoc=8

**\*\*Deliverables:\*\***

- `custom\_sweep.py` script
- Results table (CSV or JSON format)
- Plot: Parameter value vs Execution Time
- Plot: Parameter value vs Hit Rate
- Brief analysis (200-300 words):
  - What trend did you observe?
  - Why does this parameter have this effect?
  - Where does performance saturate?

### **### Part 3: Multi-Parameter Analysis (40 minutes)**

**\*\*Objective:\*\*** Understand parameter interactions

**\*\*Tasks:\*\***

1. Run the full parameter sweep:

```
``` python scripts/cache_sweep.py \
--gem5=../../build/RISCV/gem5.opt \
--config=configs/cache_config.py \
--binary=<path_to_binary> \
--output=full_sweep_results
```

2. Analyze results using:

```
``` python scripts/analyze_sweep.py \
full_sweep_results/results.json \
--output=analysis_plots
```

```

3. Examine the generated plots and statistics

****Deliverables:****

- `results.json` from complete sweep
- Analysis plots (at minimum):
 - L1D size impact
 - L2 size impact
 - Associativity impact
- Summary statistics (mean, min, max for key metrics)
- Top 3 configurations ranked by:
 - Lowest execution time
 - Highest L1D hit rate
 - Highest L2 hit rate

Part 4: Design Analysis & Recommendations (30 minutes)

****Objective:**** Make data-driven microarchitecture design decisions

****Tasks:****

1. Compare your single-parameter results with full sweep results
2. Identify the "Pareto-optimal" configurations:
 - Configs where you can't improve one metric without hurting another

3. Answer the following questions:

****a) Performance Bottlenecks (5 points)****

- Is execution time dominated by L1D misses, L2 misses, or memory stalls?
- What percentage of memory requests reach main memory?

****b) Cache Efficiency (5 points)****

- Which cache level has the best hit rate? Why?
- Is L2 size or associativity more important?

****c) Cost-Benefit Trade-off (5 points)****

- What's the smallest L1D+L2 configuration that achieves 90% of peak performance?
- How much performance do you lose by using direct-mapped caches (assoc=1)?

****d) Design Recommendations (5 points)****

- Recommend an optimal configuration for:
 - **Power-constrained system** (minimize cache size)
 - **High-performance system** (maximize performance)
 - **Balanced system** (best performance/cost ratio)
- Justify your recommendations with data

****Deliverables:****

- Answers to all 4 questions (500-800 words total)
- Graphs comparing Pareto-optimal vs suboptimal configs
- Design recommendation table with justification

File Organization

assignment_cache_optimization/

```
|—— benchmarks/  
|   |—— matrix_multiply.c  
|—— configs/  
|   |—— cache_config.py  
|—— scripts/  
|   |—— cache_sweep.py
```

```
|   └── analyze_sweep.py  
|   ├── results/  
|   |   ├── single_param_sweep/  
|   |   ├── full_sweep_results/  
|   |   └── analysis_plots/  
└── ASSIGNMENT-1.docx (this file)
```

Compilation

To compile the matrix multiply benchmark for RISCV:

```
```# Requires RISCV cross-compiler  
riscv64-unknown-linux-gnu-gcc -O2 -static benchmarks/matrix_multiply.c -o matrix_multiply
...``
```

### ### Optional Enhancements

- Test with different matrix sizes (64x64, 256x256)
- Vary memory bandwidth
- Compare different CPU types (TimingSimpleCPU vs O3CPU)
- Add power estimation metrics

### ## Submission Checklist

**Note: Everything in a single zipped directory (“problem\_1\_assignment\_1\_soln.zip”) to be uploaded on Moodle (one submission per group).**

- [ ] Part 1: Test run output and screenshots
- [ ] Part 2: Custom sweep script, results table, 2 plots, 200-300 word analysis
- [ ] Part 3: Full sweep results.json, 3+ analysis plots, statistics summary
- [ ] Part 4: Answers to all 4 design questions (500-800 words), recommendation table
- [ ] README with instructions to reproduce results
- [ ] All code properly commented and documented

### ## Additional Resources

- [gem5 Documentation](<https://www.gem5.org/documentation/>)